

PushAround: Collaborative Path Clearing via Physics-Informed Hybrid Search

Abstract—The passage of large vehicles in cluttered environments is often blocked by movable obstacles, motivating the use of mobile robot teams to proactively clear traversable corridors. Existing approaches to Navigation Among Movable Obstacles (NAMO) mainly plan abstract sequences of obstacle displacements but often neglect physical feasibility, overlooking robot dimensions, obstacle mass, contact geometry, and required forces, which can lead to strategies that fail in execution. This work introduces *PushAround*, a physics-informed hybrid search framework that incrementally constructs executable clearing plans by coupling geometric connectivity reasoning with short-horizon simulation. During search, repeated W-clearance connectivity graph (WCCG) queries expose frontier gaps where the vehicle corridor must be widened and rank them using estimated downstream gap-sequence costs. Guided by the selected gap, the search generates direct or recursive pushing candidates by choosing which obstacles to displace, in which directions, and with which contact points and forces; only physics-validated candidates are then admitted into the clearing plan. This interleaved geometry–physics screening avoids exhaustive rollout over obstacle–action combinations while preserving physically grounded plan expansion, as demonstrated through extensive simulations and hardware experiments.

I. INTRODUCTION

In many real-world settings, the passage of large vehicles or transport units is obstructed by movable obstacles such as pallets, boxes, or equipment, motivating the deployment of mobile robot teams to actively clear traversable corridors and enable safe passage. This capability is critical in cluttered and unstructured environments such as warehouses, disaster sites, and dense urban spaces. Conventional path planning methods assume static obstacles and thus fail once direct routes are blocked [1]. Navigation Among Movable Obstacles (NAMO) has been studied as a means of creating new routes by pushing, pulling, or rotating obstacles [2], yet most methods abstract away physical feasibility. Factors such as robot dimensions, obstacle size and mass, contact geometry, applied forces, and the coupled dynamics of pushing are typically ignored, producing plans that are *geometrically valid but physically unrealizable*. As shown in Fig. 1, this difficulty is amplified when a large robot cannot reach obstacle boundaries, while a small robot lacks sufficient pushing force. Cooperative pushing by several small robots is a viable alternative, but introduces strong physical coupling since simultaneous pushes can trigger chained motions of multiple obstacles, making prediction and planning substantially more complex.

A. Related Work

NAMO extends motion planning to cluttered environments by planning obstacle displacements through pushing, pulling, or rotating [2], [3]. Later work improved scalability using



Fig. 1. Collaborative path clearing in both simulation (**Top**) and hardware (**Bottom**). Two robots actively push movable obstacles to create a W-clear path for a large vehicle to traverse from start to goal, via a hybrid search over the sequence of obstacles to push and the associated pushing modes.

heuristic search, learning-based local planning, and search-based planning in interactive environments [3], [4], [5]. Recent studies further consider movability-aware pushing with physics-based rollouts [6] and collaborative multi-robot pushing [7]. However, many NAMO formulations still simplify robot geometry, actuation limits, object mass, and contact dynamics, leading to physically unrealizable plans.

By employing multiple robots, collaborative pushing provides a potential solution for such scenarios and has been widely studied in multi-robot manipulation. Foundational studies examined the mechanics of pushing and the limit-surface model [8], [9]. Subsequent work considered cooperative transport with robot swarms [10], multi-robot box pushing with reinforcement learning [11], collaborative pushing of polytopic objects in cluttered scenes [7], and learning-based loco-manipulation for long-horizon quadrupedal pushing [12]. These studies demonstrate effective multi-robot cooperation, but typically focus on a prescribed object and avoid exploiting inter-object contacts for multiple objects.

Physics-informed planning has emerged to incorporate dynamics models, contact reasoning, or simulation feedback into motion generation. Learned dynamics have been used for humanoid contact planning [13], while multi-contact force-control methods address physically consistent execution [14]. For non-prehensile manipulation, demonstration-guided optimal control reasons over contact modes and pushing dynamics [15], and physics-based simulation is interleaved with multi-agent pathfinding for manipulation among movable objects [16], [17]. Neural motion-planning methods further use learned physical priors for motion generation [18], [19]. While these approaches improve physical grounding, they either remain focused on mobile manipulator, contact execution, or generic planning, and do not directly solve collaborative path clearing with many interacting movable obstacles.




Fig. 2. Overview of the proposed framework. **Top:** Scene inputs, WCCG construction, gap ranking, and sampling of pushing modes for breaking frontier gaps; **Bottom-left:** Sim-in-the-loop hybrid search with parallel validation in simulation; **Bottom-right:** Online execution and adaptation during path clearing.

B. Our Method

As shown in Fig. 2, this work introduces **PushAround**, a physics-informed framework for collaborative multi-robot pushing that actively constructs traversable corridors for large vehicles in cluttered environments. The central novelty is a hybrid search that jointly determines which obstacles to displace and how to push them, coupling the high-level sequence of obstacle clearing with the low-level pushing modes that specify contact points, directions, and forces. The framework integrates three components: a W -clearance connectivity graph (WCCG) that certifies vehicle passage under width constraints, a gap-ranking strategy that prioritizes frontier gaps by estimated effort, and the pushing mode generation under geometric constraints. These components are encapsulated under a search procedure that incrementally expands candidate clearing plans while validating their physical realizability. This integration yields plans that are not only geometrically valid but also feasible under robot dimensions, object mass, contact geometry, and coupled pushing dynamics. Efficiency is achieved by combining compact geometric reasoning with prioritized evaluation. Extensive simulations and experiments are conducted for validation.

The contributions are twofold: (I) it introduces a unified framework that couples collaborative pushing with physics-informed verification, producing executable plans for path clearing in cluttered environments; and (II) it shows significant improvements in feasibility, efficiency and scalability over existing NAMO and collaborative clearing methods.

II. PROBLEM DESCRIPTION

A. Workspace and Robots

The workspace is a bounded planar region $\mathcal{W} \subset \mathbb{R}^2$ that contains two types of obstacles. A set $\mathcal{O} \subset \mathcal{W}$ represents immovable structures, while a set of $M \geq 0$ movable rigid polygons is defined as $\Omega \triangleq \{\Omega_1, \dots, \Omega_M\} \subset \mathcal{W}$. Each movable obstacle Ω_m is characterized by its mass M_m , inertia I_m , frictional parameters, and state $\mathbf{s}_m(t) \triangleq (\mathbf{x}_m(t), \psi_m(t))$, where $\mathbf{x}_m(t) \in \mathbb{R}^2$ is the planar position of its centroid and $\psi_m(t) \in \mathbb{R}$ its orientation. The region occupied by Ω_m at time t is denoted $\Omega_m(t)$. Moreover, a small team of N robots, indexed as $\mathcal{R} \triangleq \{R_1, \dots, R_N\}$, operates as a cooperative unit. Each robot R_i is modeled as a rigid body with state

$\mathbf{s}_{R_i}(t) \triangleq (\mathbf{x}_{R_i}(t), \psi_{R_i}(t))$, where $\mathbf{x}_{R_i}(t) \in \mathbb{R}^2$ is its position and $\psi_{R_i}(t) \in \mathbb{R}$ its orientation. The footprint of R_i is denoted $R_i(t)$. The instantaneous free space is given by:

$$\widehat{\mathcal{W}}(t) \triangleq \mathcal{W} \setminus \left(\mathcal{O} \cup \{\Omega_m(t)\}_{m=1}^M \cup \{R_i(t)\}_{i=1}^N \right), \quad (1)$$

where $\widehat{\mathcal{W}}(t)$ excludes regions occupied by immovable obstacles, movable obstacles, and robots.

B. External Vehicle and Clearance Goal

An external vehicle $\mathcal{V}$ whose footprint is enclosed by a disk of radius $W/2 > 0$ must navigate within the workspace from a start position $\mathbf{x}_V^S \in \mathcal{W}$ to a goal position $\mathbf{x}_V^G \in \mathcal{W}$. Since movable obstacles may obstruct the way, the clearance goal is to rearrange them so that the workspace contains a passage from $\mathbf{x}_V^S$ to $\mathbf{x}_V^G$ that can accommodate this enclosing disk, i.e., a passage of width at least W . At time t , this is expressed by the admissible reference-center space $\mathcal{F}_W(t) \triangleq \{p \in \widehat{\mathcal{W}}(t) \mid \mathbb{B}_{W/2}(p) \subset \widehat{\mathcal{W}}(t)\}$, where $\mathbb{B}_{W/2}(p)$ is the disk of radius $W/2$ centered at p . The W -clearance condition is then given by:

$$\exists \mathcal{P}_V^W \subset \mathcal{F}_W(t) : \mathbf{x}_V^S \rightsquigarrow \mathbf{x}_V^G, \quad (2)$$

which certifies a collision-free centerline path for any vehicle footprint contained in the enclosing disk.

C. Collaborative Pushing Modes

To enlarge blocked gaps, the robots may actively reconfigure Ω by pushing movable obstacles. For obstacle $\Omega_m \in \Omega$, a pushing mode is denoted by $\xi_m \triangleq (\mathbf{c}_m, \mathbf{u}_m)$, where $\mathbf{c}_m = (c_1, \dots, c_N)$ is an ordered tuple of contact points on $\partial\Omega_m$, and the ordering specifies the robot-to-contact assignment. The vector $\mathbf{u}_m \in \mathbb{R}^{2N}$ stacks the corresponding planar contact forces, expressed in the body frame of Ω_m . For a contact tuple $\mathbf{c}_m$, the set of admissible forces $U_m(\mathbf{c}_m) \subset \mathbb{R}^{2N}$ collects contact forces satisfying the robot pushing limits, robot-object friction, and unilateral contact constraints. The body wrench induced by a mode is defined as:

$$\mathbf{q}_m(\xi_m) \triangleq G_m(\mathbf{c}_m)\mathbf{u}_m, \quad (3)$$

where $G_m(\mathbf{c}_m) \in \mathbb{R}^{3 \times 2N}$ is the contact-to-wrench mapping; and $\mathbf{q}_m(\xi_m) \in \mathbb{R}^3$. The admissible mode set is state dependent as denoted by $\Xi_m(\mathbf{s})$. It contains modes whose contacts are




Fig. 3. Illustration of the W-clearance connectivity graph (WCCG) and gap ranking. From left to right: cluttered PyBullet scenario, WCCG overlay with centroid nodes, bridge nodes, and annotated gap widths, induced WCCG faces, and the prioritized frontier-gap sequence.

reachable from the current robot configurations, pre-contact configurations are feasible in the current clutter, and forces satisfy $\mathbf{u}_m \in U_m(\mathbf{c}_m)$. Hence, applying a feasible mode $\xi_m \in \Xi_m(\mathbf{s}(t))$ yields the physical transition:

$$\mathbf{s}(t^+) \triangleq \Phi(\mathbf{s}(t), m, \xi_m), \quad (4)$$

where $\mathbf{s}(t)$ stacks all robot and obstacle states.

D. Problem Statement

The goal is to compute a *hybrid schedule* for the robots that reconfigures the movable obstacles so that the vehicle admits a W -feasible path from start to goal. The overall schedule for the robotic fleet is defined as: $\pi \triangleq \{(\Omega_{m_k}, \xi_k, \Delta t_k)\}_{k=1}^K$, where $\Omega_{m_k} \in \Omega$ specifies the movable obstacle to manipulate; $\xi_k \in \Xi_{m_k}$ is the pushing mode; and $\Delta t_k > 0$ is the duration of this mode. Thus, the optimization balances the execution time and the physical effort of the clearance process, i.e.,

$$\min_{\pi} \left\{ T + \alpha \sum_{k=1}^K J(m_k, \xi_k, \mathbf{s}(t_k)) \right\}, \quad (5)$$

where $T \triangleq \sum_{k=1}^K \Delta t_k$ is the total task duration; $\alpha > 0$ is a trade-off parameter; and $J(\cdot)$ is a simulation-based control effort or feasibility cost evaluated at state $\mathbf{s}(t_k)$. Note that the above optimization problem is constrained by (1)–(4), which jointly ensure collision-free evolution within the workspace, valid dynamic transitions under pushing, and a terminal W -clearance path for the vehicle.

Remark 1. Problem above uniquely couples obstacle selection, pushing mode optimization and switching into one hybrid optimization. This yields a combinatorial search space far more complex than classical NAMO, with simple objects or hand-crafted contact modes [8], [10], [11]. ■

III. PROPOSED SOLUTION

As shown in Fig. 2, PushAround first uses the W-Clearance Connectivity Graph in Sec. III-A and gap ranking in Sec. III-B to identify and prioritize blocking gaps. It then generates and evaluates physically feasible collaborative pushing modes in Sec. III-C. These components are integrated into the simulation-in-the-loop hybrid search in Sec. III-D, followed by the online execution procedure summarized in Sec. III-E.

A. W-Clearance Connectivity Graph

The feasibility of routing the vehicle reduces to whether a W -width corridor connects $\mathbf{x}_V^S$ and $\mathbf{x}_V^G$. A W -clearance connectivity graph (WCCG) is introduced here as shown in Fig. 3,

which encodes obstacle adjacencies under clearance W and enables *fast connectivity queries*. Constructed directly from geometry rather than grids or roadmaps, it avoids discretization error and remains consistent, which is crucial as connectivity is checked repeatedly during planning.

1) *Graph Construction*: The WCCG is built by decomposing all clearance-relevant obstacle and workspace boundaries into convex components or segments. From each component C , a centroid node v_c is created. When two components C_u and C_v have closest points $p_u \in C_u$ and $p_v \in C_v$ with distance less than W , bridge nodes are added at p_u and p_v . These bridge nodes are connected by a bridge-bridge edge, annotated with the gap width $w_{uv} \triangleq \|p_u - p_v\|$, and further linked back to their centroids with centroid-bridge edges. Narrow passages with $w_{uv} < W$ are marked as potential bottlenecks. The resulting graph is given by:

$$\mathcal{G}_W \triangleq (\mathcal{V}, \mathcal{E}_c \cup \mathcal{E}_b), \quad (6)$$

where $\mathcal{V}$ are centroid and bridge nodes; $\mathcal{E}_c$ contains centroid-bridge edges; and $\mathcal{E}_b$ for bridge-bridge edges with widths w_{uv} .

2) *Connectivity Criterion*: Once $\mathcal{G}_W$ has been constructed, connectivity queries can be performed without explicitly computing a geometric path. Let $\text{Face}(\mathcal{G}_W)$ denote the set of connected planar faces induced by the embedded WCCG, i.e., the maximal regions separated by obstacle components and WCCG bottleneck edges. The criterion for (2) is given by:

$$(\mathbf{x}_V^S, \mathbf{x}_V^G \in \mathcal{F}_W(t)) \text{ and } (\exists F \in \text{Face}(\mathcal{G}_W) : \mathbf{x}_V^S, \mathbf{x}_V^G \in F), \quad (7)$$

which requires the start and goal to be individually W -clear and to lie in the same WCCG face. This is sufficient for the graph-level query because bridge-bridge edges with width below W encode non-traversable bottlenecks, while points in the same face are not separated by such a bottleneck. A frontier-tracing procedure, similar to the BugPlanner [20], is then used to extract a skeleton Σ as an ordered sequence of centroid and bridge nodes. Starting from $\mathbf{x}_V^S$, it casts a ray toward $\mathbf{x}_V^G$ and follows the encountered frontier loop until either a valid exit is found or a blocking cycle is returned.

Remark 2. Unlike sampling or grid-based planners, the WCCG avoids workspace discretization and thus eliminates resolution-induced errors. Its purely geometric construction also enables efficient connectivity and clearance queries, which are repeatedly invoked by the hybrid planner below. ■

B. Ranking of Frontier Gaps

When the condition in (7) fails, the vehicle cannot reach $\mathbf{x}_V^G$ from $\mathbf{x}_V^S$ through a W -clear path $\mathcal{P}_V^W$. In this case, the planner




Fig. 4. **Direct** pushing moves a gap-side blocker along a gap-opening direction, whereas **recursive** pushing first clears movable blockers in the swept region and then returns to the root gap-opening task.

must prioritize *blocking gaps* on the reachable frontier of $\mathcal{G}_W$ according to their predicted downstream cost. As shown in Fig. 3, the ranked list provides the local clearance targets that guide the hybrid search in the sequel.

1) *Frontier Extraction and Candidate Gaps*: A BugPlanner query on $\mathcal{G}_W$ above either confirms connectivity or returns a counter-clockwise frontier loop $\mathcal{L}$ that separates the two endpoints. The bridge-bridge edges on $\mathcal{L}$ define the first-hop candidate set $\Gamma_{\mathcal{L}} \triangleq \{g_1, \dots, g_K\}$, where each gap g_k is represented by its two closest boundary points, its width $w_{g_k} < W$, and the adjacent obstacles that can potentially be moved. These candidates are then ranked by predicted downstream cost and passed to the hybrid search as local clearance targets.

2) *Evaluation for Immediate Cost*: Each candidate $g \in \Gamma_{\mathcal{L}}$ is evaluated by combining the cost for the robots to reach the gap and the effort required to widen it. Let $s_{\mathcal{R}}$ be the current robot positions, and $\mathbf{o}_g$ as the outside insertion point of gap g . The resulting one-hop cost is given by:

$$C(g | \mathcal{L}, s_{\mathcal{R}}) = \lambda_t C_t(s_{\mathcal{R}}, \mathbf{o}_g) + \lambda_p (C_{dp}(g) + C_{rp}(g)) \quad (8)$$

where $\lambda_t, \lambda_p > 0$ are weighting factors for the transition and pushing costs, respectively; function $C_t(\cdot)$ denotes the collision-free distance between two points; the function $C_{dp}(\cdot)$ measures the widening effort of direct pushing, which increases with narrower opening and heavier adjacent obstacles; the function $C_{rp}(\cdot)$ estimates the pushing effort if recursive pushing is required, i.e., multiple out-layer obstacles should be cleared *before* g can be widened.

3) *Long-term Cost w.r.t. Goal*: Furthermore, to prioritize gaps closer to the goal, a heuristic is added to the evaluation, i.e., $h(g) = \eta \|\mathbf{o}(g) - \mathbf{x}_V^G\|$, with the scaling constant $\eta > 0$. Lastly, a short A*-style presearch is adopted to cross each candidate gap, recompute the next frontier, and continue for a limited width and depth. The *cost-to-connect* is predicted by:

$$\widehat{\text{Cost}}(g) \triangleq C(g | \mathcal{L}, s_{\mathcal{R}}) + \sum_{g' \in \Pi^*(g)} \{C(g' | \cdot) + h(g')\}, \quad (9)$$

where $\Pi^*(g)$ is the *sequence of subsequent gaps* discovered after virtually crossing g ; the dots indicate updated inputs along the rollout of consecutive pushes; and the scalar score $\widehat{\text{Cost}}(g) > 0$ for each candidate gap. Consequently, the final output of this module is the ranked list of candidate gaps, i.e., $\text{Rank}(\Gamma_{\mathcal{L}}) \triangleq \text{argsort}_{g \in \Gamma_{\mathcal{L}}} \widehat{\text{Cost}}(g)$, in ascending predicted




Fig. 5. Reachable contacts are filtered by the desired object twist, assigned to robots, refined by local replacement, and finally validated by simulation.

cost. This ordered set is passed to the hybrid search module, such that only the promising gaps are expanded first.

Remark 3. Efficiency of the above ranking procedure can be improved by caching frontier loops, storing the transitions $(\mathcal{L}, g) \mapsto \mathcal{L}'$, and accelerating the edge queries with axis-aligned bounding-box culling. ■

C. Pushing Mode Generation and Evaluation

For a selected frontier gap g , this subsection constructs a set of candidate pushing strategies $\Xi_g \triangleq \{(\Omega_m, \mathbf{v}, \xi)\}$, where $\mathbf{v} \triangleq (v_x, v_y, \omega) \in \mathbb{R}^3$ is a short-horizon pushing direction for the movable obstacle Ω_m induced by g , and $\xi \triangleq (c_m, \mathbf{u})$ is the associated pushing mode in (4).

1) *Direct and Recursive Pushing Tasks*: For *direct push*, let Ω_m denote one of the two movable obstacles adjacent to g . For each such obstacle, a small set of candidate body velocities $\{\mathbf{v}_m^k\}_{k=1}^{K_m}$ is sampled with a bias toward enlarging the gap g , while a few angular components are included to account for local geometry. However, if the swept region of a direct push intersects movable blockers, *recursive push* tasks are further required. Namely, for each out-layer blocker Ω_m , the nominal direction moves it away from both the swept region and the reference gap sequence $\Pi^*(g)$, with the candidate velocities $\{\mathbf{v}_m^k\}_{k=1}^{K_m}$. The recursive pushes aim to move the blocker out of the swept region rather than directly open the gap. Recursion therefore creates a bounded dependency chain which prioritizes the final blocker Ω_m , as shown in Fig. 4.

2) *Pushing Mode Generation*: At a search state s , a pushing mode is generated for each candidate direction $\mathbf{v}_m^k$ and obstacle Ω_m . In this subsection, contact tuples, object velocities, forces, and wrenches are expressed in the body frame of Ω_m . For a contact tuple $\mathbf{c}$, its multi-directional force feasibility is:

$$J_{\text{MF}}^m(\mathbf{c}, \mathbf{v}_m^k) \triangleq \sum_{\mathbf{v}_d \in \mathcal{D}_m^k} w_d \min_{\mathbf{u} \in U_m(\mathbf{c})} \|\mathbf{q}_m(\mathbf{c}, \mathbf{u}) + \mathbf{q}_m^\mu(\mathbf{v}_d)\|, \quad (10)$$

where $\|\cdot\|$ is a weighted ℓ_1 wrench norm, $\mathcal{D}_m^k$ is a finite perturbation set around $\mathbf{v}_m^k$, and $\mathbf{q}_m^\mu(\mathbf{v}_d)$ is the quasi-static friction wrench of Ω_m ; the inner minimization is a small Linear Program (LP) over the admissible contact forces $U_m(\mathbf{c})$. A state- and direction-dependent contact set $\mathcal{A}_m^k(s) \subset \partial\Omega_m$ is obtained from the robot-reachable boundary of Ω_m by




Fig. 6. **Top:** Core components of the proposed physics-informed hybrid search; **Bottom:** Sampled push tasks are evaluated via parallel simulation, producing updated states for search expansion.

velocity–force compatibility filtering, as illustrated in Fig. 5. Starting from cached layouts, if available, and random initializations in $\mathcal{A}_m^k(\mathbf{s})$, the contact tuple is refined by sequential greedy replacement. At the ℓ -th greedy step, let $\mathbf{c}^\ell$ be the current tuple, and let $\mathbf{c}_{r,c}^\ell$ be obtained by replacing its r -th contact with $c \in \mathbb{R}^2$. The best replacement is derived by:

$$\mathbf{c}^* = \underset{c \in \mathcal{A}_m^k(\mathbf{s})}{\operatorname{argmin}} J_{\text{MF}}^m(\mathbf{c}_{r,c}^\ell, \mathbf{v}_m^k), \quad (11)$$

where the greedy refinement for $\mathbf{c}^*$ sweeps over contact slots until J_{MF}^m no longer decreases. Among initializations, the best local tuple is selected as $\mathbf{c}_m^{k,*}$, thus $\mathbf{u}_m^k$ is obtained by the inner minimization in (10) for $\mathbf{v}_m^k$. This yields $\xi_m^k = (\mathbf{c}_m^{k,*}, \mathbf{u}_m^k)$. The mode is inserted into a persistent MODE-CACHE, indexed by discretized shape, mass, available robot force, and pushing direction, and used only as a warm start for later similar queries. The best few pairs $(\mathbf{v}_m^k, \xi_m^k)$ with $\xi_m^k \in \Xi_m(\mathbf{s})$ are retained for gap g and passed to simulation evaluation.

3) *Pushing Strategy Evaluation:* For a current state $\mathbf{s}$ and a candidate strategy $(\Omega_m, \mathbf{v}, \xi) \in \Xi_g$, a short-horizon rollout evaluates the transition in (4) under the selected mode. As shown in Fig. 5, a strategy is accepted only if the robots reach the pre-contact poses, maintain stable contact, and induce obstacle motion above threshold. The simulation returns:

$$(\mathbf{s}', \delta_T, \delta_J) \triangleq \text{EvalSim}(\mathbf{s}, (\Omega_m, \mathbf{v}, \xi)), \quad (12)$$

where $\mathbf{s}'$ is the successor state given the strategy; $\delta_T > 0$ is the time cost; and δ_J is the control effort. For efficiency, the pushing rollout may be bypassed when the mode has been validated before and the swept region of Ω_m is collision-free.

D. Physics-Informed Hybrid Search

The proposed physics-informed hybrid search (PIHS) selects blocking gaps and validates the corresponding multi-robot pushing strategies through short-horizon simulation, as shown in Fig. 6 and summarized in Alg. 1. Unlike purely geometric NAMO planners, PIHS expands only physics-validated strategies. Candidate strategies are evaluated in parallel.

1) *Tree and Initialization:* The search tree $\mathcal{T}$ is composed of nodes $\nu \triangleq (\mathbf{s}, \pi)$, where $\mathbf{s}$ is the current system state of all robots and movables in (4), and π the partial pushing strategy in (5). Two global functions are maintained: $\text{Rank}(\nu)$

Algorithm 1: Physics-Informed Hybrid Search

Input: $\mathbf{s}_0, \alpha, \text{EvalSim}(\cdot), W$ -criterion by (7)

Output: π^*

```

1  $\nu_0 \leftarrow (\mathbf{s}_0, \emptyset), \chi(\nu_0) = 0;$ 
2 Init  $\mathcal{Q}$  by (13);
3 while not terminated do
4    $\nu \leftarrow \mathcal{Q}.\text{pop\_min}()$  by (13);
5   foreach  $g \in \text{Rank}(\nu)$  do
6     Compute  $\Pi^*(g)$ , generate  $\Xi_g$ ;
7     foreach  $(\Omega_m, \mathbf{v}, \xi) \in \Xi_g$  do
8        $\tau = (\Omega_m, \mathbf{v}, \xi);$ 
9       // Direct or recur.
10       $(\mathbf{s}', \delta_T, \delta_J) \leftarrow \text{EvalSim}(\nu, \tau)$  by (12);
11      if success then
12         $\nu' \leftarrow (\mathbf{s}', \pi \cup \{\tau\});$ 
13         $\chi(\nu') \leftarrow \chi(\nu) + \delta_T + \alpha \delta_J;$ 
14        Add  $\nu'$  to  $\mathcal{Q}$ ;
15      if  $\mathbf{s}'$  admits  $\mathcal{P}_V^W$  by (7) then
16         $\pi^* \leftarrow \pi',$  break;
17 return  $\pi^*;$ 

```

stores ranked candidate gaps $\Gamma_{\mathcal{L}_\nu}$ given the frontier loop $\mathcal{L}_\nu$ associated with node ν , together with their exploration costs by (9); and $\chi(\nu)$ returns the cumulative execution cost from the root to ν . The root is initialized as $\nu_0 \triangleq (\mathbf{s}_0, \emptyset)$, where $\mathbf{s}_0$ is the initial system state.

2) *Node Selection:* At each iteration, the node with minimum best-first priority is selected from the queue. Priority balances realized cost and estimated remaining effort:


$$f(\nu) \triangleq \chi(\nu) + \min_{g \in \text{Rank}(\nu)} \widehat{\text{Cost}}(g), \quad (13)$$

where $\chi(\nu)$ is the cumulative execution cost so far, and the second term is the minimum predicted remaining cost among the unexplored gaps of ν . This scoring ensures that nodes are expanded in an order that jointly accounts for physical effort already incurred and the most promising future actions.

3) *Node Expansion with Parallel Simulation:* Let the selected node be $\nu \triangleq (\mathbf{s}, \pi)$. For each ranked gap $g \in \text{Rank}(\nu)$, its predicted gap sequence $\Pi^*(g)$ in (9) is first derived. Then, the mode optimization by (10) provides a candidate strategy set $\Xi_g \triangleq \{(\Omega_m, \mathbf{v}, \xi)\}$, including both the direct push of the adjacent movable obstacles, and recursive push of out-layer obstacles. Thus, expansion is prioritized first by the gap ordering in $\text{Rank}(\nu)$, and then by the candidate ordering within Ξ_g . Each candidate strategy $(\Omega_m, \mathbf{v}, \xi) \in \Xi_g$ is first checked by a geometric quick-pass and, if it passes, evaluated by the short-horizon rollout in (12). A successful evaluation yields $(\mathbf{s}', \delta_T, \delta_J)$, from which a child node $\nu' \triangleq (\mathbf{s}', \pi')$ is created by appending $\tau \triangleq (\Omega_m, \mathbf{v}, \xi)$ to the current partial schedule π . The cumulative cost is updated by: $\chi(\nu') \triangleq \chi(\nu) + \widehat{C}(\nu, \tau)$, where $\widehat{C}(\nu, \tau) \triangleq \delta_T + \alpha \delta_J$ is the realized incremental cost. Note that all candidate strategies in the current batch are evaluated *concurrently in parallel*, and a single expansion can produce multiple validated successors.

4) *Termination:* Termination occurs when a node ν reaches a state $\mathbf{s}$ for which a W -clear path $\mathcal{P}_V^W$ exists from $\mathbf{x}_V^s$ to $\mathbf{x}_V^G$

Scenario 1: 10 movable obstacles, 2 robots



Scenario 2: 14 movable obstacles, 2 robots




Fig. 7. Execution and planning alignment in Scenario 1 (Left) and Scenario 2 (Right). **Top:** simulated snapshots during path clearing. **Bottom:** corresponding WCCG overlays with the start face in blue. The executed pushes progressively expand the reachable W -clearance region until the goal is connected.

by (7). Thus, the schedule π stored in ν constitutes a complete solution to (5), encoding both the sequence of gaps to clear and the pushing strategies that realize them.

Remark 4. Efficiency comes from three lightweight heuristics: gap-lookahead reduces unproductive branches, MODE-CACHE warm-starts mode optimization, and the quick-pass rollout bypass skips simple motions whose mode was previously validated, avoiding redundant simulations. ■

E. Overall Analyses

1) *Execution and Online Adaptation:* The hybrid search returns a pushing schedule $\pi = \{\tau_1, \dots, \tau_K\}$, where $\tau_k \triangleq (\Omega_k, v_k, \xi_k)$ specifies the target obstacle, pushing direction and mode. The schedule is executed sequentially in a receding-horizon manner. For each τ_k , the robots move to the required mode ξ_k for obstacle Ω_k and apply the pushing direction v_k from the current state as in (12). During transition, robots are reassigned by preserving the circular order of contact points along the obstacle boundary [7], which reduces path crossings under sufficient clearance. After each push, the system state is updated from observation, and execution proceeds to τ_{k+1} if the action succeeds. For consistency with execution, the WCCG is rebuilt periodically from the current configuration. Execution terminates once a W -clear path $\mathcal{P}_V^W$ connects $\mathbf{x}_V^S$ and $\mathbf{x}_V^G$. Replanning is triggered when contact reassignment is blocked, the realized motion deviates from simulation, a push fails, or the updated geometry invalidates the remaining schedule. In such cases, Alg. 1 is called from the updated state.

2) *Computational Complexity:* Given M movable obstacles, let $\mathcal{G}_W$ be the current WCCG in (6), Γ_ν be the visible frontier-gap set at node ν , and let Ξ_g be the candidate pushing-strategy set for gap g . Then one node expansion in Alg. 1 consists of: (I) WCCG update and frontier extraction, with complexity $\mathcal{O}(M \log M + |\mathcal{V}| + |\mathcal{E}_c| + |\mathcal{E}_b|)$; (II) gap ranking over $\Gamma_{\mathcal{L}_\nu}$, with complexity $\mathcal{O}(|\Gamma_{\mathcal{L}_\nu}| \log |\Gamma_{\mathcal{L}_\nu}|)$; and (III) candidate generation and validation over the ranked gaps. If C_{LP} denotes the cost of one LP in mode generation, T_{sim} the rollout horizon of one $\text{EvalSim}(\cdot)$ in (12), n_w the number of parallel simulations, and $\rho_\nu \in [0, 1]$ the fraction of validated candidates, then the dominant complexity is $\mathcal{O}(\sum_{g \in \text{Rank}(\nu)} |\Xi_g| (C_{LP} + \frac{\rho_\nu T_{sim}}{n_w}))$. Hence, the overall cost per expansion is dominated by the simulation process, while quick-pass, early stopping and parallel rollouts can reduce the constant factors in practice.

IV. EXPERIMENTS

Numerical experiments are conducted to validate the proposed PIHS algorithm. The algorithms are implemented in Python3 and PyBullet [21] on a computer with an Intel Core i9-14900KF and an NVIDIA RTX 4080. Simulation and hardware videos are provided in the supplementary material.

A. Numerical Simulations

1) *Setup and Configuration:* The simulation uses a step of $\Delta t = 1/240$ s and a control period of $1/40$ s. Robots are modeled as disk or box pushers with sizes consistent with the W -clearance definition. Movable obstacles have masses uniformly sampled from $[5, 15]$ kg, while immovables are modeled with zero mass. A trial is successful when a W -clear path exists from start to goal and the vehicle reaches its goal disc. The nominal scenario consists of an 8×5 m bounded workspace with two bar obstacles forming bottlenecks and a mixed pool of curved and polygonal shapes, including rings, ellipses, X/T/L/diamond, arrow-like shapes, rectangles, and cylinders. Two robots start in the lower-left and the goal lies on the right, with 10–15 movables placed randomly under a minimum separation. The per-task simulation horizon is 80 steps, and up to 64 candidate push tasks are generated per expansion. The node priority follows (13) with heuristic factor 10, and gap sampling follows a softmax distribution with temperature 0.05. Unless otherwise stated, quantitative comparisons and ablations are conducted on Scenario 1 with 10 movable obstacles over 40 randomized trials, while Scenario 2 with 14 movable obstacles is used as an additional representative run in a more cluttered layout.

2) *Results:* Two representative simulation runs are shown in Fig. 7. Scenario 1 contains 10 movable obstacles and two robots. In the initial state, the start and goal belong to different W -clearance faces, indicating that the target passage is blocked by movable obstacles. Guided by the selected frontier gaps, the planner obtains a clearing plan with 6 push tasks in 7.4 s. During execution, the planned pushes progressively enlarge the reachable W -clearance region. The execution finishes in 42.27 s, and the final WCCG overlay shows that the start face contains the goal. Meanwhile, Scenario 2 increases to 14 movable obstacles. The planner finds a 9-task clearing plan in 11.2 s, and the execution finishes in 61.27 s. Compared with Scenario 1, the longer task sequence reflects the additional bottlenecks and stronger obstacle interactions. In both scenarios, the updated WCCG overlays are aligned

TABLE I
PERFORMANCE COMPARISON AND ABLATIONS ON SCENARIO 1
OVER 40 RANDOMIZED TRIALS.

Method / Variant	Succ.(%)	PT (s)	ET (s)	#Sims	#Pushes
DFS-WCCG	25.0	>100.0	51.2	>500	8.0
SL-Push (offline)	62.5	0.06	44.1	0.0	7.3
SL-Push (sim)	75.0	18.2	64.6	20.0	10.2
Rec-NAMO	37.5	13.3	42.4	0.0	7.8
PushAround (ours)	97.5	10.3	28.6	121.5	6.0
w/o gap-lookahead	82.5	16.9	35.2	178.3	7.2
w/o mode-cache	92.5	17.1	33.4	164.6	6.5
w/o quick-pass	90.0	23.6	41.9	208.2	6.1
w/o reinsertion	85.0	12.8	31.1	151.7	7.8



Fig. 8. Failure cases by baselines Rec-NAMO, SL-Push and DFS-WCCG.

with the executed pushes: the blue start face expands after each clearance until a complete W -clear path is found.

3) *Comparisons*: Four literature-inspired baselines are evaluated on Scenario 1 over 40 randomized trials under the same W -clearance criterion, object distribution, and execution simulator. The baselines include: **DFS-WCCG**, a simulation-in-the-loop depth-first search guided by the WCCG with fixed axis-aligned pushes, following the spirit of search-based NAMO and Bug-style frontier navigation [3], [20]; **SL-Push (offline)**, a geometric route-level pushing baseline that clears blockers by minimal displacements along a straight or waypointed route without physics validation, related to local planning among pushable or movable objects [2], [4]; **SL-Push (sim)**, which validates the same route-level push candidates through physics simulation, in line with simulation-based planning for non-prehensile manipulation among movable objects [16]; and **Rec-NAMO**, a simulation-free recursive NAMO baseline that combines Dijkstra routing with local blocker recursion, following geometric and search-based NAMO formulations [5]. Table I shows that PushAround achieves the highest success rate of 97.5% and the lowest combined planning and execution time of 38.9 s. DFS-WCCG suffers from exponential branching, requiring more than 100 s of planning. SL-Push (offline) plans almost instantly in 0.06 s, but its non-physical displacements often become unrealizable during execution, resulting in 62.5% success. SL-Push (sim) improves success to 75.0% by validating pushes in simulation, but increases the planning time to 18.2 s and requires 10.2 pushes on average. Rec-NAMO often fails to complete a connected W -clear corridor in dense clutter.

4) *Ablation Studies*: Moreover, four internal ablations are conducted: **w/o gap-lookahead**, which ranks frontier gaps without downstream costs; **w/o mode-cache**, which disables warm starts from previously validated modes; **w/o quick-**




Fig. 9. Scalability in dense clutter with 45 movable objects. **Top**: snapshots during the path clearing of 24.5 s. **Bottom**: WCCG overlays.




Fig. 10. Density scaling in the 8×8 m workspace. The panels report success rate, mean planning time, and mean execution time.

pass, which evaluates all candidate strategies by rollout; and **w/o reinsertion**, which prevents promising but temporarily failed nodes from being reconsidered. Table I shows that gap-lookahead and reinsertion mainly improve robustness, with success dropping more than 10% when either is removed. Removing mode-cache causes a smaller success drop to 92.5%, consistent with its role as a warm start rather than a feasibility gate. Quick-pass mainly improves efficiency: disabling it increases the average number of simulations from 121.5 to 208.2 and the planning time from 10.3 s to 23.6 s.


5) *Density Scalability*: Density scalability is evaluated in an 8×8 m workspace with $N \in \{15, 20, 25, 30, 35, 40, 45\}$ movable obstacles and five seeds per density, as shown in Fig. 9. For each seed, denser instances are generated by incrementally adding objects. Unsuccessful trials, including plan failures and timeout-pruned rows, are assigned the 400 s cap in planning-time statistics. As shown in Fig. 10, PushAround maintains 100% success up to $N = 45$, while its capped mean planning time increases only from 2.6 s to 49.2 s. In contrast, Rec-NAMO fails completely from $N = 35$, and SL-Push and DFS-WCCG reach only 20% success at $N = 45$. The non-recursive ablation further drops to 40% success at $N = 45$, confirming that recursive clearance is critical for the dense scenes.

B. Hardware Experiments

1) *System Description*: Experiments are conducted in a $5 \text{ m} \times 6 \text{ m}$ workspace assembled from interlocking $0.6 \text{ m} \times 0.6 \text{ m}$ foam mats, as illustrated in Fig. 11. Each trial employs 2 UGVs and 6 movable obstacles represented by cardboard boxes covered with colored paper for visual distinction. Both robots and movables are equipped with three to four motion-capture markers. A motion-capture system streams global poses to PyBullet in real time, enabling policy execution and visualization. Movables are randomly placed for each trial.

2) *Results*: Two representative real-world runs are shown in Fig. 11. In Scenario 1, PIHS breaks five candidate gates

Scenario 1



Scenario 2




Fig. 11. Real-world pushing experiments with the execution-planning alignment (over 20 runs and 2 scenarios). **Top**: snapshots of two robots pushing obstacles in a 5×6 m workspace with movable objects and fixed boundaries. **Bottom**: WCCG overlays with the start face in blue.

within 6.5 s, using 12 node expansions, 49 visited system-state snapshots, and a maximum search depth of 6. The four-task schedule first pushes and rotates the yellow L-shaped object to open a gap near an immovable obstacle. It then pushes the green and orange rectangular objects in opposite lateral directions, progressively expanding the reachable W -clearance region. After the green cylinder is pushed to the right, a complete W -clear path is established at $t = 52$ s. In Scenario 2, the proposed method succeeds in 8/10 trials by generating an average of 5.2 pushes. The planning phase takes 11.4 s on average, and first selects a direct push on the blue L-shaped object, which also moves the adjacent T-shaped block and opens the first gap into the narrow passage. The complete execution takes 77.1 s on average with 15.9 s deviations. The main deviations arise from direction-dependent force limits, drifting, and Mecanum-wheel slippage near 45° motion directions, which are compensated by online replanning.

V. CONCLUSION

This work presents a multi-robot framework PushAround for path clearing in unstructured environments using a hybrid search algorithm that jointly plans obstacles selection, contact points, and pushing forces efficiently. The method demonstrates robustness to execution deviations and hardware uncertainty. Future directions include obstacle states estimation without external sensing, uncertain masses, and limited robot visibility for real-world applications.

REFERENCES

- [1] L. Liu, X. Wang, X. Yang, H. Liu, J. Li, and P. Wang, "Path planning techniques for mobile robots: Review and prospect," *Expert Systems with Applications*, vol. 227, p. 120254, 2023.
- [2] M. Stilman and J. J. Kuffner, "Navigation among movable obstacles: Real-time reasoning in complex environments," *International Journal of Humanoid Robotics*, vol. 2, no. 04, pp. 479–503, 2005.
- [3] M. Stilman, J.-U. Schamburek, J. Kuffner, and T. Asfour, "Manipulation planning among movable obstacles," in *Proceedings 2007 IEEE international conference on robotics and automation*, 2007, pp. 3327–3332.
- [4] L. Yao, V. Modugno, A. M. Delfaki, Y. Liu, D. Stoyanov, and D. Kanoulas, "Local path planning among pushable objects based on reinforcement learning," in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 2024, pp. 3062–3068.
- [5] Z. Ren, B. Suvonov, G. Chen, B. He, Y. Liao, C. Fermuller, and J. Zhang, "Search-based path planning in interactive environments among movable obstacles," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2025, pp. 533–539.
- [6] J. J. Weeda, S. Bakker, G. Chen, and J. Alonso-Mora, "Pushing through clutter with movability awareness of blocking obstacles," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2025, pp. 512–518.
- [7] Z. Tang, Y. Zhang, and M. Guo, "Pushingbots: Collaborative pushing via neural accelerated combinatorial hybrid optimization," *IEEE Transactions on Robotics*, vol. 42, pp. 579–599, 2026.
- [8] S. Goyal, A. Ruina, and J. Papadopoulos, "Limit surface and moment function descriptions of planar sliding," in *IEEE International Conference on Robotics and Automation (ICRA)*, 1989, pp. 794–795.
- [9] K. M. Lynch, H. Maekawa, and K. Taniguchi, "Manipulation and active sensing by pushing using tactile feedback," in *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, vol. 1, 1992, pp. 416–421.
- [10] J. Chen, M. Gauci, W. Li, A. Kolling, and R. Groß, "Occlusion-based cooperative transport with a swarm of miniature mobile robots," *IEEE Transactions on Robotics*, vol. 31, no. 2, pp. 307–321, 2015.
- [11] Y. Wang and C. W. De Silva, "Multi-robot box-pushing: Single-agent q-learning vs. team q-learning," in *IEEE/RSJ international conference on intelligent robots and systems*, 2006, pp. 3694–3699.
- [12] Y. Feng, C. Hong, Y. Niu, S. Liu, Y. Yang, and D. Zhao, "Learning multi-agent loco-manipulation for long-horizon quadrupedal pushing," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2025, pp. 14441–14448.
- [13] Y.-C. Lin, B. Ponton, L. Righetti, and D. Berenson, "Efficient humanoid contact planning using learned centroidal dynamics prediction," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2019, pp. 5280–5286.
- [14] Q. Rouxel, S. Ivaldi, and J.-B. Mouret, "Multi-contact whole-body force control for position-controlled robots," *IEEE Robotics and Automation Letters*, vol. 9, no. 6, pp. 5639–5646, 2024.
- [15] T. Xue, H. Girgin, T. S. Lembono, and S. Calinon, "Demonstration-guided optimal control for long-term non-prehensile planar manipulation," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2023, pp. 4999–5005.
- [16] D. M. Saxena and M. Likhachev, "Planning for complex non-prehensile manipulation among movable objects by interleaving multi-agent pathfinding and physics-based simulation," in *IEEE International Conference on Robotics and Automation (ICRA)*, 2023, pp. 8141–8147.
- [17] K. Ren, G. Wang, A. S. Morgan, L. E. Kavvaki, and K. Hang, "Object-centric kinodynamic planning for nonprehensile robot rearrangement manipulation," *IEEE Transactions on Robotics*, vol. 41, pp. 5761–5780, 2025.
- [18] R. Ni and A. H. Qureshi, "Progressive learning for physics-informed neural motion planning," in *Robotics: Science and Systems*, 2023.
- [19] Y. Liu, R. Ni, and A. H. Qureshi, "Physics-informed neural mapping and motion planning in unknown environments," *IEEE Transactions on Robotics*, vol. 41, pp. 2200–2212, 2025.
- [20] K. N. McGuire, G. C. H. E. de Croon, and K. Tuyt, "A comparative study of bug algorithms for robot navigation," *Robotics and Autonomous Systems*, vol. 121, p. 103261, 2019.
- [21] E. Coumans and Y. Bai, "Pybullet, a python module for physics simulation for games, robotics and machine learning," <http://pybullet.org>, 2016–2019.